

Non-life — Assignment NL3

Niels Keizer* and Robert Jan Sopers†

October 3, 2016

1 GLMs and the Lee-Carter mortality model

Q1

Assuming $\beta_\Sigma := \sum \beta_x \neq 0$ and the transformation

$$\alpha'_x \leftarrow \alpha_x + \beta_x \bar{\kappa}; \quad \kappa''_t \leftarrow \kappa_t \beta_\Sigma; \quad \beta'_x \leftarrow \frac{\beta_x}{\beta_\Sigma}; \quad \kappa'_t \leftarrow \kappa_t'' - \bar{\kappa}'' \quad \forall x, t \quad (1)$$

a straightforward substitution of this transformation shows that

$$\begin{aligned} \mu'_{xt} &= \exp(\alpha'_x + \beta'_x \kappa'_t) \\ &= \exp\left(\alpha_x + \beta_x \bar{\kappa} + \frac{\beta_x}{\beta_\Sigma}(\kappa_t'' - \bar{\kappa}'')\right) \\ &= \exp\left(\alpha_x + \beta_x \bar{\kappa} + \frac{\beta_x}{\beta_\Sigma}(\kappa_t \beta_\Sigma - \bar{\kappa} \beta_\Sigma)\right) \\ &= \exp(\alpha_x + \beta_x \bar{\kappa} + \beta_x \kappa_t - \beta_x \bar{\kappa}) = \exp(\alpha_x + \beta_x \kappa_t) = \mu_{xt} \end{aligned}$$

and together with $\bar{\kappa} = \frac{1}{T} \sum_{t=1}^T \kappa_t$ we have (assuming $\beta_\Sigma \neq 0$)

$$\sum_x \beta'_x = \sum_x \frac{\beta_x}{\beta_\Sigma} = \frac{1}{\beta_\Sigma} \sum_x \beta_x = \frac{1}{\beta_\Sigma} \beta_\Sigma = 1 \quad (2)$$

and

$$\sum_{t=1}^T \kappa'_t = \sum_{t=1}^T (\kappa_t'' - \bar{\kappa}'') = \sum_{t=1}^T \kappa_t \beta_\Sigma - \sum_{t=1}^T \beta_\Sigma \bar{\kappa} = T \bar{\kappa} \beta_\Sigma - T \beta_\Sigma \bar{\kappa} = 0 \quad (3)$$

which are the desired properties for β'_x and κ'_t .

Q2

From (3) from the assignment we have that

$$\log(\mu_{xt}) = \alpha_x + \beta_x \kappa_t \Rightarrow \alpha_x = \log(\mu_{xt}) - \beta_x \kappa_t \quad (4)$$

assuming now the property $\sum_t \kappa_t = 0$ we see that summing over t gives

$$\sum_{t=1}^T \alpha_x = T \alpha_x = \sum_{t=1}^T \log \mu_{xt} - \beta_x \sum_{t=1}^T \kappa_t = \sum_{t=1}^T \log \mu_{xt} \quad (5)$$

*Student number: 10910492

†Student number: 0629049

which implies

$$\alpha_x = \frac{1}{T} \sum_{t=1}^T \log \mu_{xt} \quad \forall x \quad (6)$$

Defining the right hand side of this equation as $\overline{\log \mu_x}$ and using the property $\sum_x \beta_x = 1$ we have by summing over x from (4) that

$$\sum_x \alpha_x = \sum_x \overline{\log \mu_x} = \sum_x \log \mu_{xt} - \kappa_t \sum_x \beta_x = \sum_x \log \mu_{xt} - \kappa_t \quad (7)$$

Solving for κ_t we find

$$\kappa_t = \sum_x (\overline{\log \mu_x} - \log \mu_{xt}) \quad \forall t \quad (8)$$

Q3

To find the α_x to minimize $\sum_{x,t} (\log m_{xt} - \alpha_x - \beta_x \kappa_t)^2$ we take the partial derivative with respect to α_x and equate to zero:

$$\frac{\partial}{\partial \alpha_x} \sum_{x',t} (\log m_{x't} - \alpha_{x'} - \beta_{x'} \kappa_t)^2 = -2 \sum_{x',t} (\log m_{x't} - \alpha_{x'} - \beta_{x'} \kappa_t) \frac{\partial \alpha_{x'}}{\partial \alpha_x} = -2 \sum_t (\log m_{xt} - \alpha_x - \beta_x \kappa_t) = 0 \quad (9)$$

where we used that $\frac{\partial \alpha_{x'}}{\partial \alpha_x} = 1$ if $x = x'$ and zero otherwise. Using $\sum_t \kappa_t = 0$ we have

$$-2 \sum_t \alpha_x = -2 \sum_t \log m_{xt} \Rightarrow \hat{\alpha}_x = \frac{1}{T} \sum_t \log m_{xt} = \overline{\log m_{xt}} \quad (10)$$

Q4

If we change the transformation to be

$$\alpha'_x \leftarrow \alpha_x + \beta_x \bar{\kappa}; \quad \kappa''_t \leftarrow \kappa_t \sqrt{\sum_x \beta_x^2}; \quad \beta'_x \leftarrow \frac{\beta_x}{\sqrt{\sum_x \beta_x^2}}; \quad \kappa'_t \leftarrow \kappa''_t - \bar{\kappa}'' \quad \forall x, t \quad (11)$$

then we still have $\mu'_{xt} = \mu_{xt}$ (the proof remains the same if we change the definition of β_Σ to $\sqrt{\sum_x \beta_x^2}$) and we have the property that

$$\sum_x (\beta'_x)^2 = \sum_x \frac{\beta_x^2}{\sum_x \beta_x^2} = 1 \quad (12)$$

The script then becomes

```
kappa.LC <- kappa.LC * sqrt(sum(beta.LC^2))
beta.LC <- beta.LC/sqrt(sum(beta.LC^2))
alpha.LC <- alpha.LC + mean(kappa.LC)*beta.LC
kappa.LC <- kappa.LC - mean(kappa.LC)
```

By construction we have $\mathbf{Z} = \mathbf{U}\Sigma\mathbf{V}'$ and as remarked we also have $\mathbf{Z}\vec{1} = \vec{0}$. This implies that

$$\vec{0} = \Sigma^{-1}\mathbf{U}'\vec{0} = \Sigma^{-1}\mathbf{U}'\mathbf{Z}\vec{1} = \Sigma^{-1}\mathbf{U}'\mathbf{U}\Sigma\mathbf{V}'\vec{1} = \mathbf{V}'\vec{1} \quad (13)$$

using $\mathbf{U}'\mathbf{U} = \mathbf{I}$ and $\Sigma^{-1}\Sigma = \mathbf{I}$. This shows that $\sum_{x,t} v_{xt} = 0$ so in particular $\sum_t v_{1t} = 0$. By definition $\hat{\kappa}_t = \sigma_1 v_{1t}$ so that

$$\sum_t \kappa_t = \sigma_1 \sum_t v_{1t} = 0 \quad (14)$$

Q5

The following relations hold for the SVD $\mathbf{Z} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}'$ of the matrix \mathbf{Z}

$$\mathbf{Z}'\mathbf{Z} = \mathbf{V}(\mathbf{\Sigma}'\mathbf{\Sigma})\mathbf{V}' \quad (15)$$

$$\mathbf{Z}\mathbf{Z}' = \mathbf{U}(\mathbf{\Sigma}\mathbf{\Sigma}')\mathbf{U}' \quad (16)$$

where the right hand sides are the eigenvalue decompositions of the left-hand sides. The columns of \mathbf{V} are the eigenvectors of $\mathbf{Z}'\mathbf{Z}$ and the columns of \mathbf{U} are then eigenvectors of $\mathbf{Z}\mathbf{Z}'$. Executing the script gives the following output

```
> u1 <- eigen(Z%*%t(Z))$vectors[,1]
> v1 <- eigen(t(Z)%*%Z)$vectors[,1]
> d1 <- sqrt(eigen(t(Z)%*%Z)$values[1])
> beta.LC1 <- u1/sum(u1)
> kappa.LC1 <- v1*d1*sum(u1)
> range(beta.LC-beta.LC1); range(kappa.LC-kappa.LC1) # 'identical'
[1] -2.298509e-17  1.075529e-16
[1] -7.105427e-14  1.563194e-13
```

which implies that $\text{beta.LC} = \text{beta.LC1}$ and $\text{kappa.LC} = \text{kappa.LC1}$ apart from round-off error. Given the relationship between the SVD and the eigendecomposition of $\mathbf{Z}'\mathbf{Z}$ respectively $\mathbf{Z}\mathbf{Z}'$ this output is expected. In u1 the first eigenvector of $\mathbf{Z}\mathbf{Z}'$ is placed which corresponds to the first vector of \mathbf{U} , in v1 the first eigenvector of $\mathbf{Z}'\mathbf{Z}$ which corresponds with \mathbf{V} and in d1 the eigenvalues of $\mathbf{Z}'\mathbf{Z}$. The assignment of beta.LC1 and kappa.LC1 is done using the transformation rules and by construction we then have that these assignments equal beta.LC respectively kappa.LC .

Q6

Running the code gives the following output

```
> library(gnm) ## install it the first time you use it
> set.seed(1)
> start <- exp(lnExt.vec + alpha.LC[x] + beta.LC[x]*kappa.LC[t])
> system.time(
+   gg <- gnm(Dxt.vec ~ 0 + offset(lnExt.vec) + x + Mult(x,t), family=poisson,
+             mustart=start, trace=TRUE)
+ ) ## ~ 13 sec
Initialising
Initial Deviance = 100722.194038
Running main iterations
Iteration 1. Deviance = 39749.368321
Iteration 2. Deviance = 25917.880223
Iteration 3. Deviance = 23870.090768
Iteration 4. Deviance = 23505.680116
Iteration 5. Deviance = 23431.747520
Iteration 6. Deviance = 23413.988074
Iteration 7. Deviance = 23408.876090
Iteration 8. Deviance = 23407.375544
```

```

Iteration 9. Deviance = 23406.912470
Iteration 10. Deviance = 23406.770225
Iteration 11. Deviance = 23406.726007
Iteration 12. Deviance = 23406.712305
Iteration 13. Deviance = 23406.708046
Iteration 14. Deviance = 23406.706724
Iteration 15. Deviance = 23406.706313
Iteration 16. Deviance = 23406.706185
Iteration 17. Deviance = 23406.706146
Iteration 18. Deviance = 23406.706133
Iteration 19. Deviance = 23406.706129
Iteration 20. Deviance = 23406.706128
Iteration 21. Deviance = 23406.706128
Iteration 22. Deviance = 23406.706128
Iteration 23. Deviance = 23406.706128
Iteration 24. Deviance = 23406.706128
Iteration 25. Deviance = 23406.706128
Iteration 26. Deviance = 23406.706128
Iteration 27. Deviance = 23406.706128
Iteration 28. Deviance = 23406.706128
Iteration 29. Deviance = 23406.706128
Iteration 30. Deviance = 23406.706128
Done
user  system elapsed
11.09    0.33    11.56
> gg$deviance; gg$iter ## 23406.706128 30
[1] 23406.71
[1] 30

```

To find the optimal parameter estimates we run `gg$coefficients` which outputs all parameter estimates. By inspection we see that the α_x parameters correspond to `gg$coefficients[1:101]`, the β_x parameters are `gg$coefficients[102:202]` (in the output `Mult(., t).x1 - Mult(., t).x101`) and the κ_t parameters are `gg$coefficients[203:260]` (in the output `Mult(x, .).t1 - Mult(x, .).t58`). The optimal parameter estimates are therefore given by

```

alpha.gnm <- gg$coefficients[1:101]
beta.gnm <- gg$coefficients[102:202]
kappa.gnm <- gg$coefficients[203:260]

kappa.gnm <- kappa.gnm * sum(beta.gnm)
beta.gnm <- beta.gnm/sum(beta.gnm)
alpha.gnm <- alpha.gnm + mean(kappa.gnm)*beta.gnm
kappa.gnm <- kappa.gnm - mean(kappa.gnm)

```

Q7

We plot the $\alpha_x, \beta_x, \kappa_t$ coefficients produced by the `gnm` and the LC in R by

```
par(mfrow=c(1,3))
```

```

plot(alpha.LC,ylim=range(alpha.LC), ylab="alpha", xlab="x", col="blue", type="l")
lines(alpha.gnm,col="red", type="l")
plot(beta.LC,ylim=range(beta.LC), ylab="beta", xlab="x", col="blue", type="l")
lines(beta.gnm,col="red", type="l")
plot(kappa.LC,ylim=range(kappa.LC), ylab="kappa", xlab="t", col="blue", type="l")
lines(kappa.gnm,col="red", type="l")

```

which gives Figure 1.




Figure 1: Coefficients $\alpha_x, \beta_x, \kappa_t$ produced by gnm (red lines) and LC (blues lines)

Q8

From equation (3) from the assignment we know that the relationship between α_x and μ_{xt} is given by $\mu_{xt} = \exp(\alpha_x + \beta_x \kappa_t)$. The increase from the log-mortality α_x (keeping all other parameters fixed) from -8 to -7 from age 16 to 18 implies that μ_{xt} increases by $\exp(1) \approx 2.72$. In the assignment about Gompertz and Makeham we have already seen this phenomenon and it is named the 'accident hump'.

Q9

We execute the code and add one line to view the coefficients of the object `g1`

```

> kappa.glm <- kappa.LC
> g1 <- glm(Dxt.vec ~ x*kappa.glm[t] + offset(lnExt.vec), poisson)

```

```

> c1 <- coef(g1)
> g1$deviance; g1$iter ## 27603.07 4
[1] 27603.07
[1] 4
> c1

```

the output given from the call `c1` is an intercept, `x2 - x101` and `kappa.glm[t], x2:kappa.glm[t] - x101:kappa.glm[t]`. We can then retrieve the optimal parameter estimates by (using the intercept resp. `kappa.glm[t]` as base contribution)

```

alpha.glm <- c(c1[1], c1[2:101] + c1[1])
beta.glm <- c(c1[102], c1[103:202] + c1[102])

```

Q10

We run the code as given in the assignment and obtain the following results

```

> g2 <- glm(Dxt.vec ~ 0 + x + t:beta.glm[x] + offset(lnExt.vec), poisson,
+          mustart=fitted(g1))
> c2 <- coef(g2)
> g2$deviance; g2$iter ## 23594.62 4
[1] 23594.62
[1] 4
> c2[c(1, nages, nages+1, nages+nyears-1, nages+nyears)] ## t58:beta.glm[x] is NA
x1          x101  t1:beta.glm[x] t57:beta.glm[x] t58:beta.glm[x]
-6.0994748    -0.5598663    88.0965362    5.6195332           NA

```

We see that the last coefficient `t58:beta.glm[x]` is NA. Using the command `summary(g2)` we see that 1 coefficient is not defined because of singularities. Therefore the last coefficient is NA. Since there is no base contribution (no constant in the regression) we can use the following to assign α_x and κ_t (adding a zero for the parameter κ_{58} because of the NA result)

```

alpha.glm <- c2[1:101]
kappa.glm <- c(c2[102:158], 0)
kappa.glm <- kappa.glm * sum(beta.glm)
beta.glm <- beta.glm / sum(beta.glm)
alpha.glm <- alpha.glm + mean(kappa.glm) * beta.glm
kappa.glm <- kappa.glm - mean(kappa.glm)

```

where the transformations ensure we have the properties of Question 1.

Q11

Using the inverse link function we can calculate `fitted(g2)[532]` in terms of `Ext.vec[532]`, `alpha.glm[x[532]]`, `beta.glm[x[532]]` and `\verb|kappa.glm[t[532]]|` as

```

exp(log(Ext.vec[532]) + alpha.glm[x[532]] + beta.glm[x[532]] * kappa.glm[t[532]])

```

We check in R if the two are equal by determining if the absolute value of the difference is small and obtain the following output

```
> abs(exp(log(Ext.vec[532])+alpha.glm[x[532]]+
beta.glm[x[532]]*kappa.glm[t[532]])-fitted(g2)[532])<0.001
x10
TRUE
```

which implies that the reconstruction and the value of `fitted(g2)[532]` are equal (checked for a maximum difference of 0.0001).

Q12

We execute the code from the assignment and obtain

```
> beta.glm <- beta.glm/sum(beta.glm)
> alpha.glm <- alpha.glm + mean(kappa.glm)*beta.glm
> kappa.glm <- kappa.glm - mean(kappa.glm)
> (d1 <- sum(dpois(Dxt.vec,Dxt.vec,log=TRUE))) ## -21643.76
[1] -21643.76
> (d2 <- sum(dpois(Dxt.vec,
+               Ext.vec*exp(alpha.glm[x] + beta.glm[x] * kappa.glm[t]),
+               log=TRUE))) ## -33441.07, same as d3
[1] -33441.07
> (d3 <- sum(dpois(Dxt.vec,fitted(g2),log=TRUE)))
[1] -33441.07
> (d4 <- log(prod(dpois(Dxt.vec,fitted(g2))))))## -Inf
[1] -Inf
> (d5 <- 2*sum(Dxt.vec*log(Dxt.vec/fitted(g2)) - (Dxt.vec-fitted(g2))))
[1] 23594.62
> (d1-d2)*2 ## 23594.62, same as d5
[1] 23594.62
```

We see that the result `d2` equals `d3` and that `d5` equals $(d1-d2)*2$. The results `d2` and `d3` are equal by construction (see Question 11). The result from `d4` should be equal to the result from `d3`. The difference is using the `log=TRUE` option in `d3` instead of taking the logarithm from the products in `d4`. The result in `d4` is not equal because R has difficulties in taking the logarithm of the very small probabilities. The equality of `d5` and $(d1-d2)*2$ follows from MART (9.29). The calculation `d5` is equal to right-hand side of MART (9.29) (with ϕ and w_i equal to zero) and $(d1-d2)*2$ is equal to the middle expression of MART (9.29).

Q13

We execute the code from the exercise and obtain the following result.

```
> range(tapply(Dxt.vec-fitted(g2),x,sum)) ## -3e-10 2e-10
[1] -3.885816e-10 2.038405e-10
```

```
> range(tapply(Dxt.vec-fitted(g2),t,sum)) ## -3300 2726
[1] -3299.962 2726.266
>
```

The first line shows the range of the differences of the data versus the fit with respect to the ages x and the second line shows the range of the differences of the data versus the fit with respect to the time t . The result shows that the residuals are very small with respect to the ages but can be significant with respect to time (residuals up between -3300 and 2726).

Q14

We run the code from the assignment and obtain the given results

```
> kappa.glm <- kappa.LC
> oldDeviance <- 0; TotnIter <- 0; start=NULL
> system.time(
+   repeat
+   { g1 <- glm(Dxt.vec~x*kappa.glm[t]+offset(lnExt.vec), poisson, mustart=start)
+     c1 <- coef(g1)
+     alpha.glm <- c(c1[1],c1[2:101]+c1[1])
+     beta.glm <- c(c1[102],c1[103:202]+c1[102])
+     g2 <- glm(Dxt.vec ~ 0+x + t:beta.glm[x] + offset(lnExt.vec), poisson,
+               mustart=fitted(g1))
+     8
+     c2 <- coef(g2)
+     alpha.glm <- c2[1:101]
+     kappa.glm <- c(c2[102:158],0)
+     kappa.glm <- kappa.glm*sum(beta.glm); beta.glm <- beta.glm/sum(beta.glm);
+     alpha.glm <- alpha.glm + mean(kappa.glm)*beta.glm
+     kappa.glm <- kappa.glm - mean(kappa.glm)
+     TotnIter <- TotnIter + g1$iter + g2$iter
+     newDeviance <- g2$deviance;
+     done <- abs((oldDeviance-newDeviance)/newDeviance)<1e-6
+     cat(g1$deviance, "\t", g2$deviance, "\n")
+     oldDeviance <- newDeviance; start <- fitted(g2)
+     if (done) break
+   }
+ ) ## ~ 6 sec
27603.07 23594.62
23416.47 23407.23
23406.74 23406.71
23406.71 23406.71
user system elapsed
4.54 0.19 4.74
> TotnIter ## 20
[1] 20
>
> AIC(g1); AIC(g2) ## 67098.22 67010.22
```



```
[1] 67098.22
[1] 67010.22
> logLik(g1); logLik(g2) ## 'log Lik.' -33347.11 with df=202 and df=158
'log Lik.' -33347.11 (df=202)
'log Lik.' -33347.11 (df=158)
```

The calculated log-likelihoods are equal for both models. Because not all parameters are determined in `g1` and `g2` the calculated AIC is not correct with respect to the complete model. Since the total effective number of parameters is $101 + 101 + 58 - 2 = 258$ (101 α_x parameters, 101 β_x parameters, 58 κ_t parameters and a redundancy of 2) the actual AIC calculation is equal to $AIC = -2l + 2k = -2(-33347.11 - 258) = 67210.22$ which was obtained from the R output

```
> -2*(logLik(g1)-258)
'log Lik.' 67210.22 (df=202)
```

Q15

The \mathbf{X} matrix has 5858 rows from the observations. For the `g1` model we estimated the α_x and β_x parameters so the matrix \mathbf{X} has 202 columns. We check the dimensions of the matrix in R and obtain the following output

```
> dim(model.matrix(g1))
[1] 5858 202
```

which are the dimensions we expected. Using the command given in the question we determine the amount of memory for each of the objects `g1` and `g2` and obtain

```
> object.size(g1); object.size(g2);
13254648 bytes
11130864 bytes
```

To determine which part of `g1` occupies the most space we use the command given and obtain the following output

```
> sort(sapply(g1,object.size))
```

rank	deviance	aic	null.deviance	iter	df.residual
48	48	48	48	48	48
df.null	converged	boundary	data	method	contrasts
48	48	48	56	96	360
control	formula	call	xlevels	terms	coefficients
544	1720	1904	5952	6840	15408
family	offset	effects	R	residuals	fitted.values
45880	46904	105952	354080	375104	375104
linear.predictors	weights	prior.weights	y	model	qr
375104	375104	375104	375104	603456	9811960

which shows that the `qr` part of the object `g1` occupies the most space.

Q16

To find the parameters \hat{b} and \hat{c} that optimize the Gompertz(b, c) likelihood we use a `glm` call with offset `offset(lnExt.vec)` and find the parameter values through exponentiating the estimates. The output of R is the following for the complete estimate, the estimate restricted to the ages 30+ and the plots.

```
> x1 <- as.numeric(x)-1
> g3 <- glm(Dxt.vec ~ x1 + offset(lnExt.vec),poisson)
> g3$coefficients
(Intercept)          x1
-9.11080428  0.08377865
>
> (b.Gompertz <- exp(coef(g3)[1]))
(Intercept)
0.0001104658
> (c.Gompertz <- exp(coef(g3)[2]))
x1
1.087388
>
> alpha.Gompertz <- log(b.Gompertz)
> kappa.Gompertz <- log(c.Gompertz)
> beta.Gompertz <- (1:nages)-1
>
> kappa.Gompertz <- kappa.Gompertz * sum(beta.Gompertz)
> beta.Gompertz <- beta.Gompertz/sum(beta.Gompertz)
> alpha.Gompertz <- alpha.Gompertz + mean(kappa.Gompertz)*beta.Gompertz
> kappa.Gompertz <- kappa.Gompertz - mean(kappa.Gompertz)
>
> g3.30 <- glm(Dxt.vec ~ x1 + offset(lnExt.vec),poisson,subset = x1>=30)
> g3.30$coefficients
(Intercept)          x1
-10.13570719  0.09782707
>
> (b.30.Gompertz <- exp(coef(g3.30)[1]))
(Intercept)
3.96386e-05
> (c.30.Gompertz <- exp(coef(g3.30)[2]))
x1
1.102772
>
> alpha.30.Gompertz <- log(b.30.Gompertz)
> kappa.30.Gompertz <- log(c.30.Gompertz)
> beta.30.Gompertz <- (1:nages)-1
> kappa.30.Gompertz <- kappa.30.Gompertz * sum(beta.30.Gompertz)
> beta.30.Gompertz <- beta.30.Gompertz/sum(beta.30.Gompertz)
> alpha.30.Gompertz <- alpha.30.Gompertz + mean(kappa.30.Gompertz)*beta.30.Gompertz
> kappa.30.Gompertz <- kappa.30.Gompertz - mean(kappa.30.Gompertz)
>
```

```

> par(mfrow=c(1,1))
> plot(alpha.LC,ylim=range(alpha.LC), ylab="alpha", xlab="x", col="blue", type="l")
> lines(alpha.Gompertz,col="black", type="l")
> lines(alpha.30.Gompertz,col="red", type="l")
>

```

The parameter estimates on the complete set are

$$b = 1.104658 \times 10^{-4} \quad \text{and} \quad c = 1.087388 \quad (17)$$

and the plot is given in Figure 2.




Figure 2: Coefficient α_x estimated by LC (blue line), Gompertz (black line) and Gompertz ages 30+ (red line)

Q17

Using the code in the assignment we obtain the Figures from the assignment.

```

g4 <- glm(Dxt.vec~x1*t-x1-1+offset(lnExt.vec), poisson, subset=x1>=30)
b <- 1e5*exp(head(coef(g4),nyears)); c <- 100*(exp(tail(coef(g4),nyears))-1)
par(mfrow=c(1,2))
plot(b, xlab="t", ylab="b*100000", ylim=c(0,10), type="l", yaxp=c(0,10,2),
main="Gompertz parameters b;\nages 30+")
plot(c, xlab="t", ylab="c-1 in %", ylim=c(9,12), type="l",
yaxp=c(9,12,3), main="Gompertz parameters c")

```

To extrapolate the trend in the graph of b_t and determine the time t for which $b_t = 0$ we use regression of b_t against t and solve for $b_t = 0$ in R. We have the following output

```
> t1 <- 1:nyears
> b.lm <- lm(b~t1,subset = t1>20)
> b.lm$coefficients
(Intercept)          t1
9.6359275   -0.1411485
>
> t.intersect <- -b.lm$coefficients[1]/b.lm$coefficients[2]
> t.intersect
(Intercept)
68.26802
>
```

At time $t = 68.27$ we have that $b_t = 0$. A negative b_t implies that there is a negative mortality rate and no people will be dying.

Q18

Using the following code:

```
b <- b/1e5; c<- c/100 + 1;
log.mortality <-function(x){ log(b) + x * log(c)}

plot(log.mortality(65), xlab="t", ylab="Log-mortality", type="l",col="black",ylim=c(-5,0))
lines(log.mortality(75), col="red", type="l")
lines(log.mortality(85), col="blue", type="l")
```

we obtain Figure 3. There is a shift of the curve upwards for higher ages and for values of t larger than 30 the log-mortality is descending with t .




Figure 3: Log-mortality for ages 65 (black line), 75 (red line) and 85 (blue line)

2 Analyzing a bonus-malus system using GLM

Q1

a)

We are asked to check if the values in Table 9.8 from MART are correct. For this, we first run the code given in the exercise:

```
> rm(list=ls(all=TRUE)) ## First remove traces of previous sessions
> fn <- "http://www1.fee.uva.nl/ke/act/people/kaas/Cars.txt"
> Cars <- read.table(fn, header=TRUE)
> Bminus1 <- Cars$B - 1; Bis14 <- as.numeric(Cars$B==14)
> Cars$A <- as.factor(Cars$A); Cars$R <- as.factor(Cars$R)
> Cars$M <- as.factor(Cars$M); Cars$U <- as.factor(Cars$U)
> Cars$B <- as.factor(Cars$B); Cars$WW <- as.factor(Cars$WW)
> ActualWt <- c(650,750,825,875,925,975,1025,1075,1175,1375,1600)
> W <- log(ActualWt/650)[Cars$WW]
>
> # GLM analysis
>
> g1 <- glm(TotCl/Expo~R+A+U+W+Bminus1+Bis14, quasipoisson, wei=Expo, data=Cars)
> g2 <- glm(TotCl/Expo~R+A+U+W+Bminus1+Bis14+M, quasipoisson, wei=Expo, data=Cars)
> g3 <- glm(TotCl/Expo~R+A+U+W+B, quasipoisson, wei=Expo, data=Cars)
>
> anova(g1,g2)
```

Analysis of Deviance Table

```
Model 1: TotCl/Expo ~ R + A + U + W + Bminus1 + Bis14
Model 2: TotCl/Expo ~ R + A + U + W + Bminus1 + Bis14 + M
  Resid. Df Resid. Dev Df Deviance
1      7515   38616941
2      7513   38614965  2    1975.8
> anova(g1,g3)
Analysis of Deviance Table
```

```
Model 1: TotCl/Expo ~ R + A + U + W + Bminus1 + Bis14
Model 2: TotCl/Expo ~ R + A + U + W + B
  Resid. Df Resid. Dev Df Deviance
1      7515   38616941
2      7504   38544506 11    72435
>
> # Multiplicative coefficients
> options(digits=7)
> exp(coef(g1)); exp(coef(g2)); exp(coef(g3))
(Intercept)          R2          R3          A2          A3          U2          W
524.3016583   1.0842682   1.1916130   0.4147224   0.6184468   1.3841303   2.3722083
  Bminus1      Bis14
  0.8978647   1.1053665
(Intercept)          R2          R3          A2          A3          U2          W
522.6627527   1.0842767   1.1914111   0.4147232   0.6184538   1.3835062   2.3721668
  Bminus1      Bis14      M2      M3
  0.8978640   1.1053568   1.0073260   1.0014581
(Intercept)          R2          R3          A2          A3          U2          W
515.5320549   1.0843018   1.1916593   0.4143437   0.6178700   1.3841612   2.3722369
  B2      B3      B4      B5      B6      B7      B8
  0.9111279  0.8275175  0.7403718  0.6842609  0.6088526  0.5416103  0.4489065
  B9      B10     B11     B12     B13     B14
  0.4151901  0.3888576  0.3459030  0.3143452  0.2832722  0.2773037
```

All coefficients can be checked individually against table 9.8 and are the same, except for models *g1* and *g2*, because the bonus malus risk factor is taken as numeric. This means that the factors in the table have been calculated from the factor for *B2* to the power $B - 1$. The coefficient for *Bminus1* differs only in the 7th decimal spot and the table is given with 4 decimals. This means that if we only need to check one of the two models. We do this by recalculating the values in R.

```
> bm_class <- seq(1,13,1)
> bm_coef <- exp((bm_class-1)*coef(g1)["Bminus1"])
> bm_coef
[1] 1.0000000 0.8978647 0.8061610 0.7238236 0.6498956 0.5835184 0.5239205
[8] 0.4704098 0.4223643 0.3792260 0.3404937 0.3057173 0.2744927
```

These values also correspond with those in table 9.8

b)

Using the coefficients of `g1`, `g2` and `g3`, compute the fitted values for the cell 4000.

For this, we use the coefficients in R. Recalculating these by hand would be rather pointless and is an exercise in working neatly over understanding the subject matter.

```
> # Observed value
> g1$y[4000]
4000
326.4545
>
> # Fitted value
> fitted(g1)[4000]; fitted(g2)[4000]; fitted(g3)[4000]
4000
634.0642
4000
636.416
4000
644.5283
```

What we can see is that all three GLM's have a fitted value that is about twice as large as the actual value. Not one of the models is close to the observed value, but the models are quite close together in their estimate.

c)

We now explain the result of the following R-code.

```
> g2$family$linkinv(model.matrix(g2)[4000,]%*%coef(g2))
      [,1]
[1,] 636.416
```

This result is equal to the fitted value of the `g2` model. This is no surprise, considering the code is equal to the definition of the fitted value for cell 4000. The inner product of the values of the risk factors and their corresponding coefficients gives the linear estimator for that cell, after which the `linkinv` function is applied, which is the exponential function. This results in the fitted value.

Q2

First we will determine the scale factor ϕ using a 'rich' model, meaning that the values of both the weight of the car and the BM class are used as factors.

```
> g.rich <- glm(TotCl/Expo~R+A+U+WW+B, quasipoisson, wei=Expo, data=Cars)
> anova(g.rich)
Analysis of Deviance Table
```

```
Model: quasipoisson, link: log
```

Response: TotCl/Expo

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev
NULL			7523	116167018
R	2	2586478	7521	113580540
A	2	23288859	7519	90291681
U	1	4479946	7518	85811735
WW	10	6931993	7508	78879742
B	13	40358336	7495	38521406

We determine the scale factor to be $\frac{38521406}{7495} = 5139.61$. Or in R:

```
> phi <- 38521406/7495
```

To check whether Bis14 can be removed from the model, we use an `anova` call on the model with Bis14 (`g1`) and without (`g.test`).

```
> g.test <- glm(TotCl/Expo~R+A+U+W+Bminus1, quasipoisson, wei=Expo, data=Cars)
```

```
> anova(g.test,g1)
```

Analysis of Deviance Table

Model 1: TotCl/Expo ~ R + A + U + W + Bminus1

Model 2: TotCl/Expo ~ R + A + U + W + Bminus1 + Bis14

	Resid. Df	Resid. Dev	Df	Deviance
1	7516	38755743		
2	7515	38616941	1	138802

Next we test if the inclusion of Bis14 is significant:

```
> test <- function (Df, Deviance){  
+   scaled.dev <- Deviance/phi  
+   test.dev <- qchisq(0.95,Df)  
+   return(scaled.dev>test.dev)  
+ }  
> test(1, 138802)  
[1] TRUE
```

First we calculate the scaled deviance. Then we calculate the 95-th percentile of the $\chi^2(k)$ distribution with Df degrees of freedom. When the improvement of scaled deviance is larger than the test value, the increase is significant. The test is implemented as a function, so it can be reused in the rest of the exercise. Also the test returns `TRUE`, therefore the inclusion of `Bis14` is a significant improvement of the model and can not be removed.

Then we check if B can be removed from model `g3`.


```

> g.test <- glm(TotCl/Expo~R+A+U+W, quasipoisson, wei=Expo, data=Cars)
> anova(g.test,g3)
Analysis of Deviance Table

Model 1: TotCl/Expo ~ R + A + U + W
Model 2: TotCl/Expo ~ R + A + U + W + B
  Resid. Df Resid. Dev Df Deviance
1      7517   78902891
2      7504   38544506 13  40358385
> test(13,40358385)
[1] TRUE

```

The test value is TRUE, so B can not be removed from the model. Next we check whether W can be removed from the model:

```

> g.test <- glm(TotCl/Expo~R+A+U+B, quasipoisson, wei=Expo, data=Cars)
> anova(g.test,g3)
Analysis of Deviance Table

Model 1: TotCl/Expo ~ R + A + U + B
Model 2: TotCl/Expo ~ R + A + U + W + B
  Resid. Df Resid. Dev Df Deviance
1      7505   45495122
2      7504   38544506  1   6950616
> test(1,6950616)
[1] TRUE

```

This result implies that W can not be removed from the model g3.

Is it helpful to allow separate coefficients for the weight class in model g1. We again check using the `anova` and `test` functions.

```

> g.test <- glm(TotCl/Expo~R+A+U+WW+Bminus1+Bis14, quasipoisson, wei=Expo, data=Cars)
> anova(g1, g.test)
Analysis of Deviance Table

Model 1: TotCl/Expo ~ R + A + U + W + Bminus1 + Bis14
Model 2: TotCl/Expo ~ R + A + U + WW + Bminus1 + Bis14
  Resid. Df Resid. Dev Df Deviance
1      7515   38616941
2      7506   38593888  9    23053
> test(9,23053)
[1] FALSE

```

This shows that allowing separate coefficients for the weight classes would not be an improvement.

Q3

To answer this question, we run the `test` function defined earlier:

```
> test(7515-7491,38616941-38408588)
[1] TRUE
```

The interaction terms do improve the model significantly. It might be worthwhile to investigate which interaction terms give the most improvement, because there might be some interaction terms which are not significant by themselves.

Q4

First we estimate the number of claims and the size per claim as described. We can combine the two models by adding their coefficients, because directly combining the two models will give a product of two exponentials, which is the same as one exponential with the arguments summed. We compare the resulting coefficients with a direct estimation.

```
> g.nCl <- glm(nCl/Expo~R+A+U+W+Bminus1+Bis14, quasipoisson, wei=Expo, data=Cars)
> g.sCl <- glm(TotCl/nCl~R+A+U+W+Bminus1+Bis14, Gamma(link="log"), wei=nCl, data=Cars)
> g.direct <- glm(TotCl/Expo~R+A+U+W+Bminus1+Bis14, quasipoisson, wei=Expo, data=Cars)
>
> mult.coef <- exp(coef(g.nCl)+coef(g.sCl))
> direct.coef <- exp(coef(g.direct))
> mult.coef; direct.coef
(Intercept)          R2          R3          A2          A3          U2
525.1107841   1.0856608   1.1901279   0.4134531   0.6145069   1.3823142
      W      Bminus1      Bis14
      2.3827096   0.8979376   1.1057074
(Intercept)          R2          R3          A2          A3          U2
524.3016583   1.0842682   1.1916130   0.4147224   0.6184468   1.3841303
      W      Bminus1      Bis14
      2.3722083   0.8978647   1.1053665
```

The resulting models have very similar results and attribute about the same amount of risk to each risk factor.